

431 Class 15

github.com/THOMASELOVE/2019-431

2019-10-15

Today's Agenda (Notes Chapters 19-20)

- ① Discussion of Quiz 1
- ② Discussion of Draft Survey for Project Study A
- ③ Statistical Inference and the dm431 data: Comparing Population Means using Independent Samples
 - Pooled t / Indicator Variable Regression
 - Welch's t
 - Wilcoxon-Mann-Whitney rank sum
 - Bootstrap with bootdif

Today's Setup and Data

```
library(magrittr); library(janitor)
library(patchwork); library(here);
library(boot); library(broom)
library(tidyverse)

source(here("R", "Love-boost.R"))

dm431 <- readRDS(here("data", "dm431.Rds"))
```

Comparing means using Independent Samples


Our population: ALL adults ages 31-70 seen for care this year and two years ago who live in Northeast Ohio with a diabetes diagnosis.

Our sample: 431 of those people, drawn in a way we hope is representative (but certainly isn't random).

- 1 Can we estimate the difference in the population mean systolic blood pressure among females in our population as compared to males in our population?
- 2 Can we estimate the difference in the population mean LDL level for those on a statin as compared to those not on a statin?
- 3 Can we estimate the difference in the population mean hemoglobin A1c for those with Medicaid vs. Medicare insurance?

SBP for females vs. males

Systolic Blood Pressure by Sex in 431 Adults with Diabetes



Numerical Summary for Two Independent Samples

```
mosaic::favstats(sbp ~ sex, data = dm431)
```

Registered S3 method overwritten by 'mosaic':

```
method          from  
fortify.SpatialPolygonsDataFrame ggplot2
```

	sex	min	Q1	median	Q3	max	mean	sd	n
1	F	90	118	128	142	208	131.1673	20.14962	257
2	M	98	120	130	140	182	131.4080	15.86577	174
	missing								
1		0							
2		0							

Confidence Interval Options for Independent Samples

- 1 Pooled t CI or Indicator Variable Regression Model (t approach assuming equal population variances)
- 2 Welch t CI (t approach without assuming equal population variances)
- 3 Wilcoxon-Mann-Whitney Rank Sum Test (non-parametric test not assuming Normality but needing symmetry to be related to means)
- 4 Bootstrap confidence interval for the difference in population means (fewest assumptions of these options)

Hypotheses Under Consideration

The hypotheses we are testing are:

- H_0 : mean in population 1 = mean in population 2 + hypothesized difference Δ_0 vs.
- H_A : mean in population 1 \neq mean in population 2 + hypothesized difference Δ_0 ,

where Δ_0 is almost always zero. An equivalent way to write this is:

- $H_0 : \mu_1 = \mu_2 + \Delta_0$ vs.
- $H_A : \mu_1 \neq \mu_2 + \Delta_0$

Yet another equally valid way to write this is:

- $H_0 : \mu_1 - \mu_2 = \Delta_0$ vs.
- $H_A : \mu_1 - \mu_2 \neq \Delta_0$,

where, again, Δ_0 is almost always zero.

Assumptions of the Pooled T test

The standard method for comparing population means based on two independent samples is based on the t distribution, and requires the following assumptions:

- 1 [Independence] The samples for the two groups are drawn independently.
- 2 [Random Samples] The samples for each of the groups are drawn at random from the populations of interest.
- 3 [Normal Population] The two populations are each Normally distributed
- 4 [Equal Variances] The population variances in the two groups being compared are the same, so we can obtain a pooled estimate of their joint variance.

The Pooled Variances t test in R

Also referred to as the t test assuming equal population variances:

```
tt <- t.test(sbp ~ sex, data = dm431, var.equal = TRUE)
tidy(tt) %>%
  select(estimate1, estimate2,
         conf.low, conf.high, method, alternative) %>%
  knitr::kable(digits = 2)
```

estimate1	estimate2	conf.low	conf.high	method	alternative
131.17	131.41	-3.82	3.34	Two Sample t-test	two.sided

Full output from Pooled T test

Two Sample t-test

```
data:  sbp by sex
t = -0.13225, df = 429, p-value = 0.8949
alternative hypothesis: true difference in means is not equal
95 percent confidence interval:
 -3.818555  3.337094
sample estimates:
mean in group F mean in group M
    131.1673      131.4080
```

Indicator Variable Regression Approach

```
model1 <- lm(sbp ~ sex, data = dm431)
tidy(model1, conf.int = TRUE, conf.level = 0.95) %>%
  select(term, estimate, conf.low, conf.high) %>%
  knitr::kable(digits = 2)
```

term	estimate	conf.low	conf.high
(Intercept)	131.17	128.89	133.44
sexM	0.24	-3.34	3.82

sexM indicator shows the effect of being Male, so the displayed CI estimates $\mu_{male} - \mu_{female}$. Invert the signs to get the $\mu_{female} - \mu_{male}$ estimate.

summary(model1)

```
> summary(model1)

Call:
lm(formula = sbp ~ sex, data = dm431)

Residuals:
    Min       1Q   Median       3Q      Max
-41.167 -11.408  -1.408   9.833  76.833

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  131.1673     1.1566  113.409  <2e-16 ***
sexM          0.2407     1.8203   0.132    0.895
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 18.54 on 429 degrees of freedom
Multiple R-squared:  4.077e-05, Adjusted R-squared:  -0.00229
F-statistic: 0.01749 on 1 and 429 DF,  p-value: 0.8949
```

Results for the SBP and Sex Study

Procedure	p for $H_0 : \mu_F = \mu_M$	95% CI for $\mu_M - \mu_F$
Pooled t	0.90	(-3.3, 3.8)

What conclusions should we draw, at $\alpha = 0.05$?

Assumptions of the Welch t test

The Welch test still requires:

- 1 [Independence] The samples for the two groups are drawn independently.
- 2 [Random Samples] The samples for each of the groups are drawn at random from the populations of interest.
- 3 [Normal Population] The two populations are each Normally distributed

But it doesn't require:

- 4 [Equal Variances] The population variances in the two groups being compared are the same.

Welch's t test is the default `t.test` in R.

Welch t test not assuming equal population variances

```
t.test(sbp ~ sex, data = dm431)
```

Welch Two Sample t-test

data: sbp by sex

t = -0.13838, df = 419.27, p-value = 0.89

alternative hypothesis: true difference in means is not equal

95 percent confidence interval:

-3.660307 3.178845

sample estimates:

mean in group F mean in group M

131.1673

131.4080

Results for the SBP and Sex Study

Procedure	p for $H_0 : \mu_F = \mu_M$	95% CI for $\mu_M - \mu_F$
Pooled t	0.90	(-3.3, 3.8)
Welch t	0.89	(-3.2, 3.7)

What conclusions should we draw, at $\alpha = 0.05$?

Assumptions of the Wilcoxon-Mann-Whitney Rank Sum Test

The Wilcoxon-Mann-Whitney Rank Sum test still requires:

- 1 [Independence] The samples for the two groups are drawn independently.
- 2 [Random Samples] The samples for each of the groups are drawn at random from the populations of interest.

But it doesn't require:

- 3 [Normal Population] The two populations are each Normally distributed
- 4 [Equal Variances] The population variances in the two groups being compared are the same.

It also doesn't really compare population means. It compares pseudo-medians again.

Wilcoxon-Mann-Whitney Rank Sum Test

```
wilcox.test(sbp ~ sex, data = dm431, conf.int = TRUE)
```

Wilcoxon rank sum test with continuity
correction

data: sbp by sex

W = 21329, p-value = 0.4167

alternative hypothesis: true location shift is not equal to 0

95 percent confidence interval:

-4.999979 2.000039

sample estimates:

difference in location

-1.99992

Results for the SBP and Sex Study

Procedure	p for $H_0 : \mu_F = \mu_M$	95% CI for $\mu_M - \mu_F$
Pooled t	0.90	(-3.3, 3.8)
Welch t	0.89	(-3.2, 3.7)

Procedure	p for $H_0 : psmed_F = psmed_M$	95% CI for M - F shift
Rank Sum	0.42	(-2.0, 5.0)

What conclusions should we draw, at $\alpha = 0.05$?

The Bootstrap

This bootstrap approach to comparing population means using two independent samples still requires:

- 1 [Independence] The samples for the two groups are drawn independently.
- 2 [Random Samples] The samples for each of the groups are drawn at random from the populations of interest.

but does not require either of the other two assumptions:

- 3 [Normal Population] The two populations are each Normally distributed
- 4 [Equal Variances] The population variances in the two groups being compared are the same.

The bootstrap procedure I use in R was adapted from Frank Harrell and colleagues. <http://biostat.mc.vanderbilt.edu/twiki/bin/view/Main/BootstrapMeansSoftware>

The bootdif function

The procedure requires the definition of a function, which I have adapted a bit, called `bootdif`, which is part of the `Love-boost.R` script we loaded earlier.

As in our previous bootstrap procedures, we are sampling (with replacement) a series of many data sets (default: 2000).

- Here, we are building bootstrap samples based on the SBP levels in the two independent samples (M vs. F).
- For each bootstrap sample, we are calculating a mean difference between the two groups (M vs. F).
- We then determine the 2.5th and 97.5th percentile of the resulting distribution of mean differences (for a 95% confidence interval).

Using the bootdif function to compare means based on independent samples

So, to compare systolic BP (our outcome) across the two levels of sex (our grouping factor) for the adult patients with diabetes in NE Ohio, run...

```
set.seed(4312019)
dm431 %>% bootdif(sbp, sex, conf.level = 0.95)
```

Mean Difference	0.025	0.975
0.2407308	-3.1197510	3.5696286

- The two columns must be separated here with a comma rather than a tilde (~), and are specified using \$ notation.
- This CI estimates $\mu_{male} - \mu_{female}$: observe the listed sample mean difference for the necessary context. Invert the signs, as before, to estimate $\mu_{female} - \mu_{male}$.

Results for the SBP and Sex Study

Procedure	p for $H_0 : \mu_F = \mu_M$	95% CI for $\mu_F - \mu_M$
Pooled t	0.90	(-3.3, 3.8)
Welch t	0.89	(-3.2, 3.7)
Bootstrap	$p > 0.05$	(-3.1, 3.6)

Procedure	p for $H_0 : psmed_F = psmed_M$	95% CI for M - F shift
Rank Sum	0.42	(-2.0, 5.0)

What conclusions should we draw, at $\alpha = 0.05$?

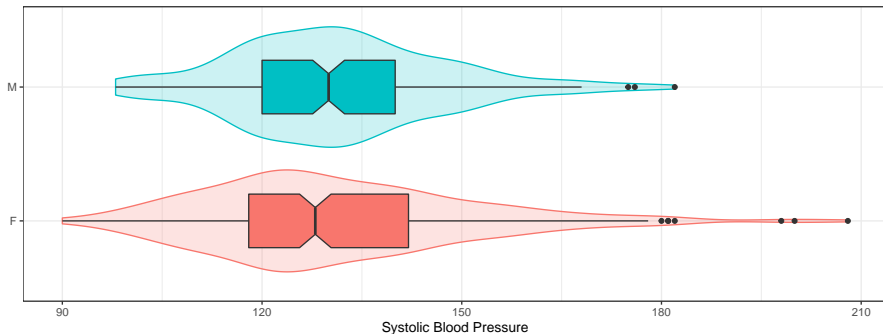
Which Method Should We Use?

- ❶ Plot the distributions of the two independent samples.
- ❷ Does it seem reasonable to assume that **each** distribution (here, both sbp in males and sbp in females) follows an approximately Normal distribution?
 - If Yes, Normal models seem appropriate, then
 - use the pooled t test (or indicator variable regression) if the sample sizes are nearly the same, or if the sample variances are quite similar
 - use the Welch's t test, otherwise (this is the default R choice)
 - If No, Normal models don't seem appropriate, then
 - compare means using the bootstrap via `bootdif`, or
 - compare pseudo-medians using the rank sum test

What did we see in our systolic BP data?

Systolic BP, within groups defined by sex


Independent Samples Comparison: SBP by Sex



LDL of statin users and non-users

Warning: Removed 37 rows containing non-finite values (stat_ydensity).

Warning: Removed 37 rows containing non-finite values (stat_boxplot).



What should we do about the missing values?

We could just look at the complete cases, and that's probably the best strategy when we're doing a two-sample t test. When we fit a more complicated regression (than just the simple indicator variable regression we need to get the pooled t test) where we're adjusting for multiple variables, then we'll consider imputation seriously.


We also probably want to turn `statin` into a factor with meaningful names.

```
dm431_sub <- dm431 %>%  
  filter(complete.cases(ldl, statin)) %>%  
  mutate(statin_f =  
    fct_recode(factor(statin),  
               statin_rx = "1", no_statin = "0"))
```

OK. Let's plot again.

LDL of statin users and non-users

Independent Samples Comparison: LDL by Statin Use



Summary Statistics for LDL by statin prescription

```
mosaic::favstats(ldl ~ statin_f, data = dm431_sub)
```

	statin_f	min	Q1	median	Q3	max	mean	sd
1	no_statin	31	76	98.0	114.5	177	97.41667	29.22364
2	statin_rx	36	70	88.5	113.0	227	96.40683	35.33276

	n	missing
1	72	0
2	322	0

Pooled t-test from indicator variable regression

```
> model_sub <- lm(ldl ~ statin, data = dm431_sub)
> summary(model_sub)

Call:
lm(formula = ldl ~ statin, data = dm431_sub)

Residuals:
    Min       1Q   Median       3Q      Max
-66.417 -25.159  -6.412  16.593 130.593

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)   97.417     4.043   24.094  <2e-16 ***
statin        -1.010     4.472   -0.226    0.821
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 34.31 on 392 degrees of freedom
Multiple R-squared:  0.00013,    Adjusted R-squared:  -0.002421
F-statistic: 0.05098 on 1 and 392 DF,  p-value: 0.8215
```

Conclusions at $\alpha = 0.05$?

Results for the LDL and Statin Study

Procedure	p for $H_0 : \mu_{statin} = \mu_{no}$	95% CI for $\mu_{statin} - \mu_{no}$
Pooled t	0.82	(-9.8, 7.8)
Welch t	0.80	(-8.9, 6.8)
Bootstrap	$p > 0.05$	(-8.7, 7.0)

Procedure	p for $H_0 : p_{smmed}_{st} = p_{smmed}_{no}$	95% CI for Statin - No shift
Rank Sum	0.27	(-12, 4)

What conclusions should we draw, at $\alpha = 0.05$?

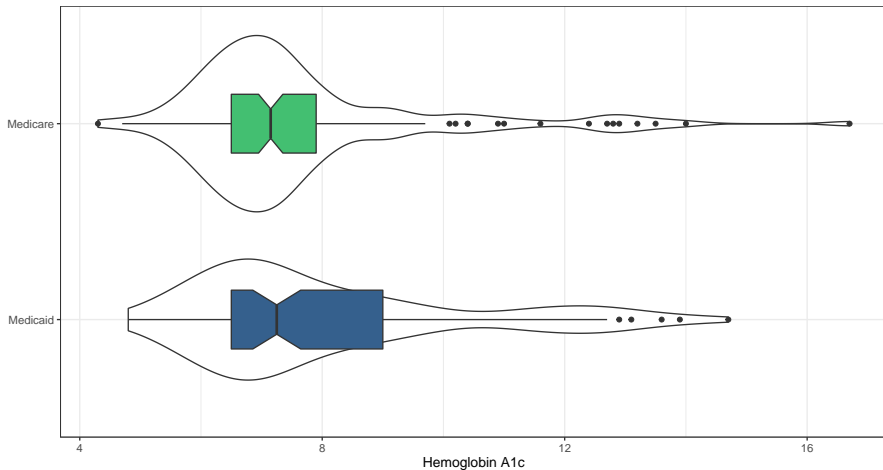
Let's compare Hemoglobin A1c for Medicare vs. Medicaid

```
dm431_sub2 <- dm431 %>%  
  filter(insurance %in% c("Medicaid", "Medicare")) %>%  
  filter(complete.cases(a1c)) %>%  
  droplevels()  
  
mosaic::favstats(a1c ~ insurance, data = dm431_sub2)
```

	insurance	min	Q1	median	Q3	max	mean	sd
1	Medicaid	4.8	6.5	7.25	9.0	14.7	8.074000	2.302946
2	Medicare	4.3	6.5	7.15	7.9	16.7	7.635246	2.041225
	n	missing						
1	100	0						
2	122	0						

Hemoglobin A1c for Medicare vs. Medicaid

Independent Samples Comparison: A1c by Insurance



Four Approaches for comparing A1c by Insurance

Procedure	p for $H_0 : \mu_{Medicaid} = \mu_{Medicare}$	95% CI for $\mu_{Medicaid} - \mu_{Medicare}$
Pooled t	0.13	(-0.14, 1.01)
Welch t	0.14	(-0.14, 1.02)
Bootstrap CI	$p > 0.05$	(-0.17, 1.02)

Procedure	p for $H_0 : p_{smcd}_{Medicaid} = p_{smcd}_{Medicare}$	95% CI for Medicaid - Medicare shift
Rank Sum	0.24	(-0.2, 0.6)

What conclusions should we draw, at $\alpha = 0.05$?

A Few Reminders About Significance

- **A significant effect is not necessarily the same thing as an interesting effect.** For example, results calculated from large samples are nearly always “significant” even when the effects are quite small in magnitude. Before doing a test, always ask if the effect is large enough to be of any practical interest. If not, why do the test?
- **A non-significant effect is not necessarily the same thing as no difference.** A large effect of real practical interest may still produce a non-significant result simply because the sample is too small.
- **There are assumptions behind all statistical inferences.** Checking assumptions is crucial to validating the inference made by any test or confidence interval.

dm431 Example 4.

Compare the proportion of current Medicare subjects whose insurance status changed to Medicare in the last two years, to the proportion of current Medicare subjects who had another insurance status two years ago?

```
dm431 %>% filter(insurance == "Medicare") %>%  
  count(insurance, insurance_old)
```

```
# A tibble: 4 x 3  
  insurance insurance_old      n  
  <fct>      <fct>      <int>  
1 Medicare Commercial     22  
2 Medicare Medicaid        4  
3 Medicare Medicare      93  
4 Medicare Uninsured        4
```

- Outcome? Exposure Groups? Paired or Independent Samples?

dm431 Example 5.

Among current Medicare subjects, compare the proportion with A1c below 8 to the proportion for the same patients two years ago.

```
dm431 %>% filter(insurance == "Medicare") %>%  
  count(a1c < 8, a1c_old < 8)
```

```
# A tibble: 7 x 3  
  `a1c < 8` `a1c_old < 8`      n  
  <lgl>      <lgl>      <int>  
1 FALSE     FALSE      16  
2 FALSE     TRUE       12  
3 FALSE     NA          1  
4 TRUE      FALSE      16  
5 TRUE      TRUE       74  
6 TRUE      NA          3  
7 NA        TRUE         1
```

- Outcome? Exposure Groups? Paired or Independent Samples?